

Deep Neural Networks for Real-time Trajectory Planning

MSc Project Presentation

Amaury Francou — *Supervisor : Dr Dante Kalise*

Imperial College London
Department of Mathematics

Presented 15/09/2022

**Imperial College
London**

Contents

- 1 Introduction
- 2 Model and optimal control problem
- 3 Generating dataset
- 4 Deep neural network for near-optimal closed loop controls
- 5 Conclusion

Introduction

NeurIPS 2017

The 2017 NeurIPS conference Pieter Abbeel main talk → presented a video in which a robot was able to perform a *wide range of household tasks* [1].



Figure 1: 2017 Neural Information Processing Systems conference (NeurIPS), Pieter Abbeel talk : *Deep learning for robotics*

- Robots have **required physical capabilities**
- But limited by a lack of **intelligent embedded controls**.

AI & Optimal Control

- Artificial intelligence and deep learning : outperforms human operators on ever-increasing range of tasks [2].
- Optimal control theory : compute **continuous controls that dynamically optimize an objective function** → can describe self-controlled robots

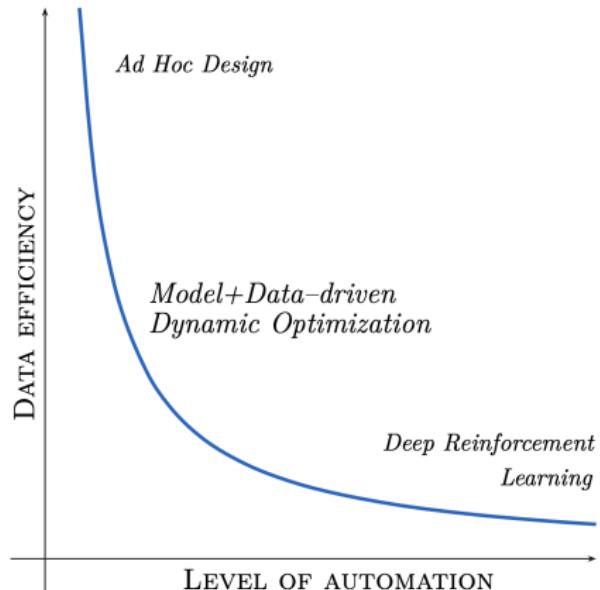


Figure 2: Automation level against data efficiency. Reprinted from [3] with permission.

Objective and Plan

- Objective : Build controller capable of navigating the agent from an initial to a final position → **trajectory planning problem**
- Criteria : Aiming for **time-energy optimality** → minimizing transportation time + magnitude of controls
- Plan :
 1. Derive physics-informed *optimal control conditions*
 2. Effortlessly generate *synthetic dataset* based on optimal trajectories
 3. Use *deep learning model* to fit the data and synthesize a controller

The quadrotor drone setting

Obtain quantitative results + leverage physics-led intuitions
→ focusing on **quadrotor drone** UAVs



Figure 3: Walkera QR X350 Quadcopter hovering [4]

- *Variety of applications such as :* infrastructure inspection - search and rescue operations - aerial photography - ...
- *Benefits from automation :* traffic management - aerial delivery - road maintenance - fire watch flights - ...

Model and optimal control problem

The planar-quadrotor

The planar-quadrotor model : 2-dimensional UAV moving in the \mathbb{R}^2 plane \rightarrow simpler dynamics **yet capturing essence of control-motion relationship.**

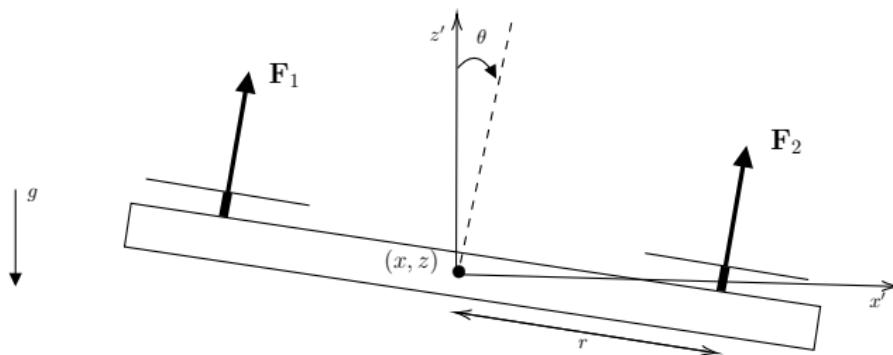


Figure 4: The planar-quadrotor model.

The model is extensively studied in the related literature [5, 6, 7, 8, 9].

Dynamics I

The dynamics governing the UAV's motion are given by Newton's laws of motion.

$$\begin{cases} m\ddot{x} = (F_1 + F_2) \sin \theta \\ m\ddot{z} = (F_1 + F_2) \cos \theta - mg \\ I\ddot{\theta} = r(F_1 - F_2) \end{cases} \quad (1)$$

Our general control is set as in the following.

$$\mathbf{u} := (u_T, u_R)^\top, \text{ where } u_T := \frac{F_1 + F_2}{m} \text{ and } u_R := \dot{\theta} \quad (2)$$

The quadrotor UAV is characterized by the *state vector* given hereby.

$$\mathbf{q} := (x, \dot{x}, z, \dot{z}, \theta)^\top \quad (3)$$

Dynamics II

The controlled dynamics are defined in the compact following way.

$$\dot{\mathbf{q}} = \mathbf{f}(\mathbf{q}, \mathbf{u}), \text{ where } \mathbf{f}(\mathbf{q}, \mathbf{u}) := \begin{pmatrix} \dot{x} \\ u_T \sin \theta \\ \dot{z} \\ u_T \cos \theta - g \\ u_R \end{pmatrix} \quad (4)$$

- Planar-quadrotor model : **arises with 5 equations dynamics**
- **The curse of dimensionality** → computing numerical optimal solutions becomes intractable as the number of equations and parameters grows
- Bypass this downfall using deep learning to compute controls **in real-time and on the edge**

Bolza problem

We consider the following *Bolza problem*.

$$\begin{aligned} \min_{\mathbf{u}(\cdot)} \quad & J[\mathbf{u}(\cdot)] := t_f + \int_0^{t_f} \|\mathbf{u}\|^2 dt \\ \text{s.t.} \quad & \mathbf{q}(t = t_f) = \mathbf{q}_f, \\ & \mathbf{q}(t = 0) = \mathbf{q}_0, \\ & \dot{\mathbf{q}} = \mathbf{f}(\mathbf{q}, \mathbf{u}), \\ & \mathbf{u} \in U, \\ & t \in [0, t_f] \end{aligned} \tag{OCP}$$

- Admissible controls in U are such that $\underline{u}_T \leq u_T(t) \leq \overline{u}_T$ and $|u_R(t)| \leq \overline{u}_R$
- Open loop controls \rightarrow solves the (OCP) equation solely based on the initial state \mathbf{q}_0 , in the form $\mathbf{u}^* = \mathbf{u}^*(t, \mathbf{q}_0)$
- **Closed loop controls** \rightarrow instant measured state is used to condition the upcoming controls, in the form $\mathbf{u}^* = \mathbf{u}^*(t, \mathbf{q})$

Pontryagin's minimum principle

- We denote λ the **costate vector** and H the **Hamiltonian**
- An optimal feedback control $\mathbf{u}^* = \mathbf{u}^*(t, \mathbf{q}, \lambda)$ verifies the following **Pontryagin's minimum principle equations**
- The (PMP) **boundary value problem provides a necessary condition for optimality**¹

$$\begin{cases} \mathbf{u}^* = \arg \min_{\mathbf{u}(\cdot) \in U} H(t, \mathbf{q}, \lambda, \mathbf{u}) \\ \dot{\mathbf{q}} = \frac{\partial H}{\partial \lambda} = \mathbf{f}(\mathbf{q}, \mathbf{u}^*), \quad \mathbf{q}(0) = \mathbf{q}_0, \quad \mathbf{q}(t_f) = \mathbf{q}_f \\ -\dot{\lambda}(t) = \frac{\partial H}{\partial \mathbf{q}}(t, \mathbf{q}, \lambda, \mathbf{u}^*(t, \mathbf{q}, \lambda)) \end{cases} \quad (\text{PMP})$$

(adjoint equation)

¹See manuscript for Hamilton-Jacobi-Bellman equation approach and relationship to PMP conditions

Solving controls

Finding optimal controls and trajectories :

1. Specifying the Hamiltonian

$$H(t, \mathbf{q}, \boldsymbol{\lambda}, \mathbf{u}) = 1 + \lambda_1 \dot{x} + \lambda_3 \dot{z} - \lambda_4 g + \left(\lambda_5 u_R + u_R^2 \right) + \left((\lambda_2 \sin \theta + \lambda_4 \cos \theta) u_T + u_T^2 \right) \quad (5)$$

2. Computing the costate variables using adjoint equation

→ The costate variables can be **parametrized by 4 constants** $c_1, c_2, c_3, c_4 \in \mathbb{R}^4$ and vary with time and state

3. Minimizing H with respect to controls

$$u_T^*(t) = \min \left(\max \left(a(t), \underline{u}_T \right), \overline{u}_T \right) \quad \text{and} \quad u_R^*(t) = \min \left(\max \left(b(t), -\overline{u}_R \right), \overline{u}_R \right) \quad (6)$$

→ The regulating function a is fully determined by the state and **the 4 constants**

→ The regulating function b is determined as an **initial value problem** parameterized by **the 4 constants**

Augmented system

- We build an **augmented system** using a vector $\mathbf{y} := (x, \dot{x}, z, \dot{z}, \theta, b)^\top$
- We solve the related augmented **initial value problem** to generate an appropriate trajectory :
 - State components are computed according to the dynamics
 - The controls u_T and u_R are computed using (6) → Using a and b

$$\dot{\mathbf{y}} = \mathbf{g}(\mathbf{y}, \mathbf{u}) := \begin{pmatrix} \dot{x} \\ u_T \sin \theta \\ \dot{z} \\ u_T \cos \theta - g \\ u_R \\ \dot{b}(t) := \left(\frac{(c_2 - c_1 t) \cos \theta - (c_4 - c_3 t) \sin \theta}{2} \right) u_T \end{pmatrix}, \quad \mathbf{y}(0) = (\mathbf{q}_0, 0)^\top \quad (7)$$

- Any choice of 4 constants c_1, c_2, c_3, c_4 fully determines a trajectory
- There exists a latent space of optimal trajectories included in \mathbb{R}^4

Generating dataset

Example trajectory I

- What is the latent space structure ?
- We generate a sizeable amount of trajectories and compare to **solver** obtained ones (ICLOCS2/IPOPT)
- We present an example starting at $\mathbf{q}_0 = (9.48, -1.12, -4.29, 7.14, 0.86)^\top$ and arriving at $\mathbf{q}(t_f = 2) = (10.64, 1.87, -3.78, -10.12, -0.87)^\top$

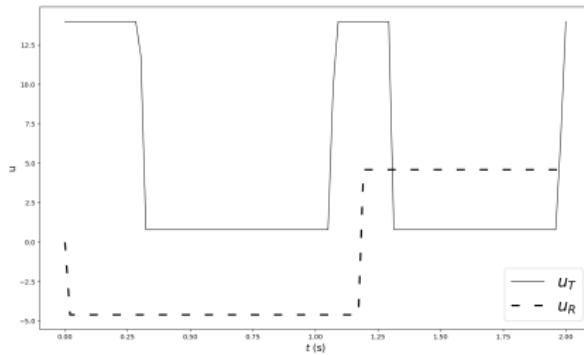
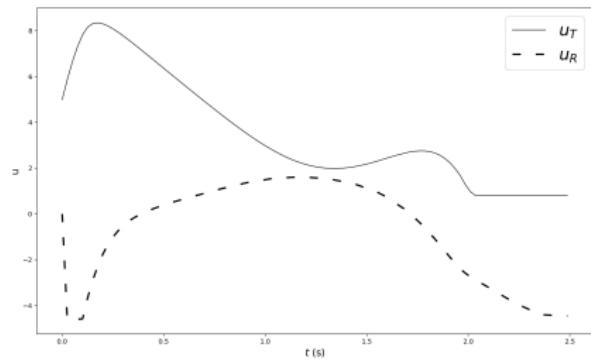


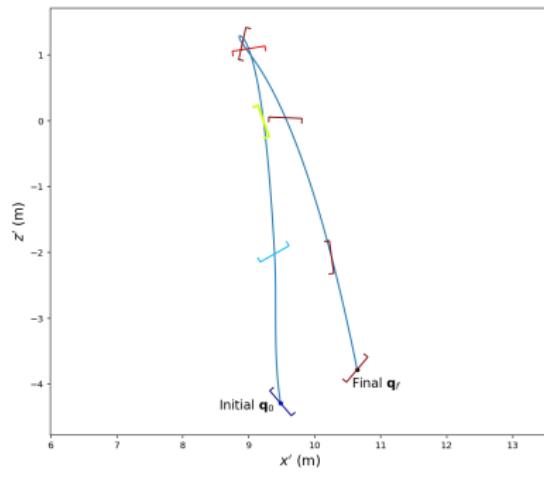
Figure 5: (a) Theoretical controls



(b) Solver-based controls

Example trajectory II

Maneuver performed in $t_f = 2$ seconds



Maneuver performed in $t_f = 2.48$ seconds

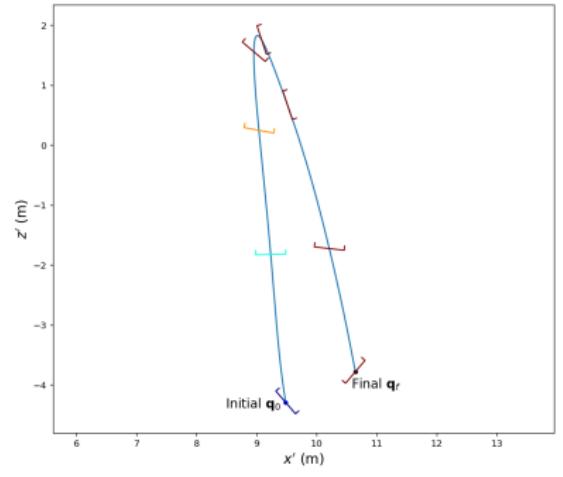


Figure 6: (a) Theoretical trajectory

(b) Solver-based trajectory

Example trajectory III

	Theoretically generated	Solver-based
t_f	2.00	2.48
E	65.42	57.36
J	67.42	59.84

Table 1: Cost functionals and contributions associated with the example trajectory.

- Theoretically generated trajectory and controls are slightly more **time-optimal**
- **Energy consumption** is lower for the solver-based controls
- The cost functional J is roughly 12.5% higher for the theoretically generated trajectory → Near-optimal
- **The latent space is a strict subset of \mathbb{R}^4**
- **Trajectories generated from any $(c_1, c_2, c_3, c_4) \in \mathbb{R}^4$ will be considered as optimal**

Data representivity

- We can **shape the data** according to our needs :
 - The constants + initial states + time horizons are **randomly sampled** (combination of Normal / log-Normal / custom uniform distributions)
 - We **require 'evenly' distributed trajectories** (relative position of initial and final locations - relative orientation - distance traveled - ...)
- **Choosing the predictor** $\hat{\mathbf{q}}(t) := \mathbf{q}_f - \mathbf{q}(t)$, with $\hat{\theta}(t) \equiv \theta_f - \theta(t) \pmod{2\pi}$ and transform the angle according to $\hat{\theta} \rightarrow (\cos \hat{\theta}, \sin \hat{\theta})$

We build the dataset $\mathcal{D} := \left\{ \left(\hat{\mathbf{q}}_{(i)}, \mathbf{u}_{(i)} \right) \right\}_{i=1}^M \subset X \times U$ using a Runge-Kutta method of order 4 implemented in Python, discarding over-represented trajectories

- Very **efficient method** → built a dataset composed of 123 750 000 points
- We use a 85%/15% split to divide the dataset into a **training set** and a **validation set**

Deep neural network for near-optimal closed loop controls

Model

The deep neural network model is composed of :

- An input layer made of 6 neurons
- 5 hidden layers made of 900 neurons each and activated using ReLU functions ($x \rightarrow \max(0, x)$)
- All connections between hidden layers are subject to dropout with a probability of 25%
- An output layer made of 2 neurons, respectively activated using **modified sigmoids** $\sigma_T(x) := (\overline{u_T} - \underline{u_T}) \frac{1}{1+e^{-x}} + \underline{u_T}$ and $\sigma_R(x) := \overline{u_R} \left(\frac{2}{1+e^{-x}} - 1 \right)$

The network is composed of 3 251 702 trainable parameters. The DNN is optimized using **stochastic gradient descent** and the training is made performing early stopping

From $z = 0$ to $z = 1$ |

- We present the 'take-off' trajectory starting at $\mathbf{q}_0 = (0, 0, 0, 0, 0)^\top$ and targeting $\mathbf{q}_f = (0, 0, 1, 0, 0)^\top \rightarrow$ The trajectory is now **closed-loop**

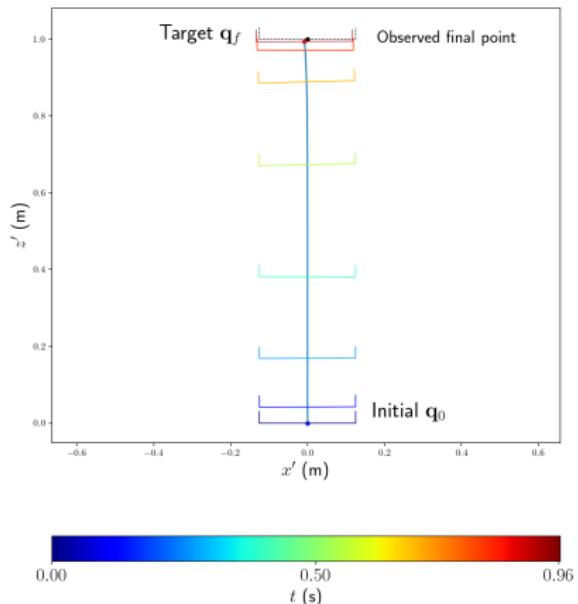


Figure 7: Closed loop quadrotor take-off trajectory.

From $z = 0$ to $z = 1$ ||

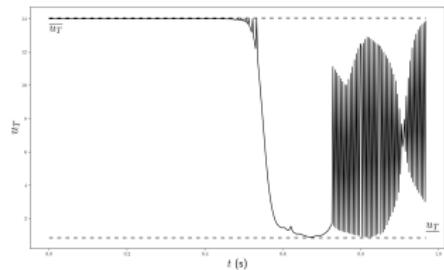
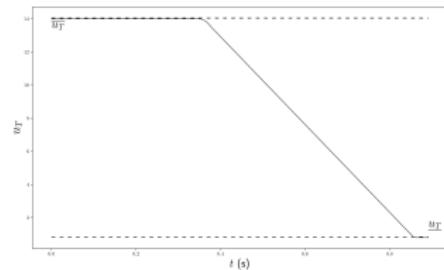


Figure 8: (a) Feedback u_T control



(b) Solver-based u_T control

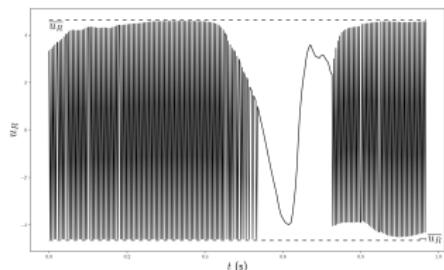


Figure 9: (a) Feedback u_R control



(b) Solver-based u_R control

Randomly selected trajectory I

- We present a randomly drawn closed loop quadrotor trajectory → starting at $\mathbf{q}_0 = (-0.69, -1.29, 4.07, 7.40, 6.21)^\top$, targeting final state $\mathbf{q}_f = (-6.51, -8.47, 10.38, 2.78, 6.25)^\top$

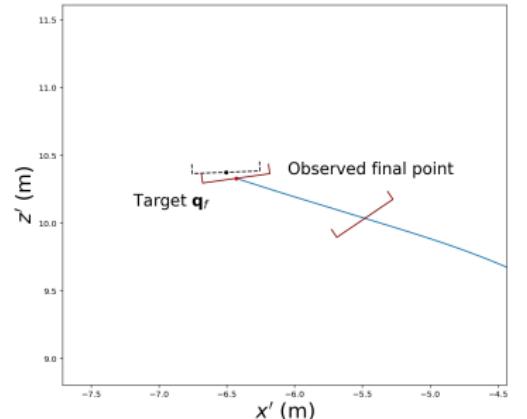
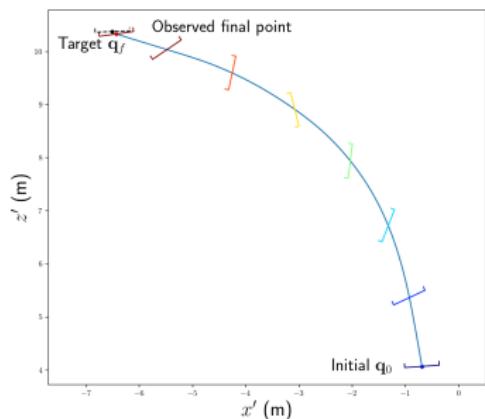


Figure 10: Randomly sampled closed loop quadrotor trajectory

Randomly selected trajectory II

- We compare the randomly sampled trajectory with solver-based one

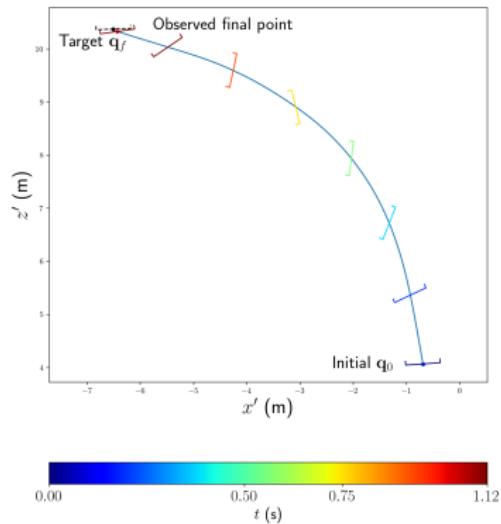
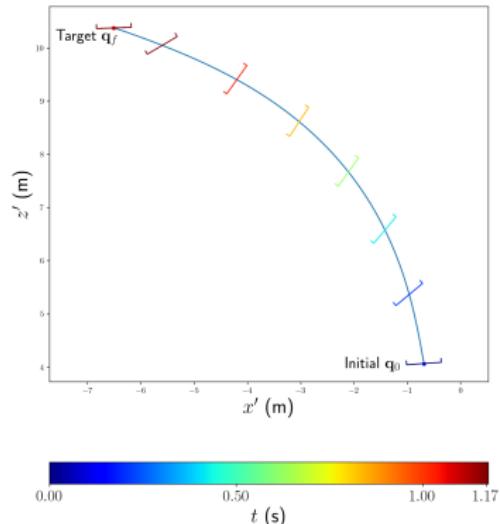


Figure 11: (a) Closed loop trajectory



(b) Solver-based trajectory

Randomly selected trajectory III

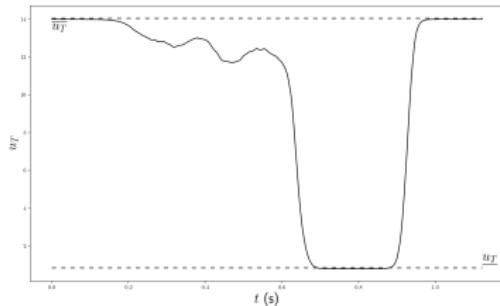
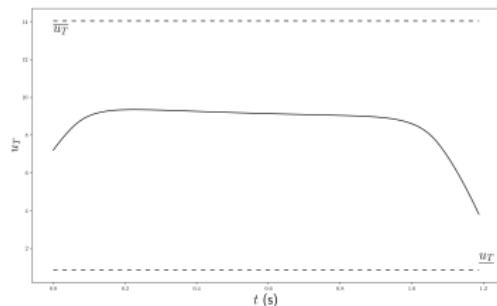


Figure 12: (a) Feedback u_T control



(b) Solver-based u_T control

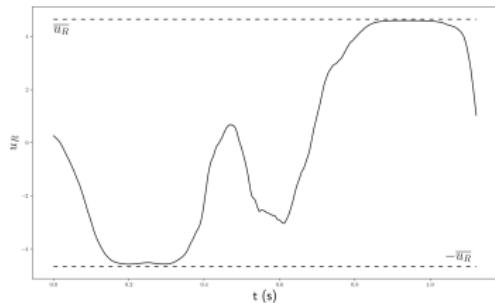
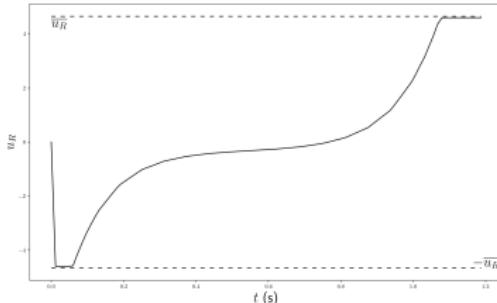


Figure 13: (a) Feedback u_R control



(b) Solver-based u_R control

Randomly selected trajectory IV

	DNN-generated	Solver-based
t_f	1.12	1.17
E	116.25	97.12
J	117.37	98.29

Table 2: Cost functionals and contributions associated with the randomly sampled trajectory.

- The closed loop controller performs the navigation in $t_f = 1.12$ seconds
- More time-optimal closed loop controls but less energy optimal → J is 19 % higher in this case
- Closed loop controls have a tendency to vary too strongly at first, before correcting in a second time
- Yet, they can be applied in real-time on the edge

Conclusion

Main results

- A sizeable dataset of near-optimal trajectories was generated efficiently
- A deep neural network was designed and trained effectively
- The quadrotor joins q_f within a reasonable tolerance autonomously for a wide range of drawn initial and final states
- Some 'atypical' final states remain un-reached by the UAV → we propose a set of improvements

Future research

Several future axis can be considered :

- Separate the controller in two different networks
- Consider augmented predictors including the time variable, in the form of $\hat{\mathbf{q}} = (x_f - x, \dot{x}_f - \dot{x}, z_f - z, \dot{z}_f - \dot{z}, \cos(\theta_f - \theta), \sin(\theta_f - \theta), t)^\top$
- Take into account progressiveness and temporal dependence of the state control relationship → Use long short-term memory (LSTM) with built-in feedback loops
- Explore latent space → A linear relationship exists between constants c_1, c_2, c_3, c_4
- Optimal control boosted deep reinforcement learning → Use deep deterministic policy gradient method (DDPG) with our kickstarted model
- Assume bang-bang controls → Go from regression to classification

References I

- [1] Pieter Abbeel. Deep learning for robotics. Neural Information Processing Systems - NeurIPS, conference, 2017.
- [2] David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dharshan Kumaran, Thore Graepel, et al. Mastering chess and shogi by self-play with a general reinforcement learning algorithm. *arXiv preprint arXiv:1712.01815*, 2017.
- [3] Dante Kalise. Week 11: Dynamic optimization, lecture notes. Optimization, MATH70005. Imperial College London, delivered 21 March 2022.
- [4] Wikimedia Commons. Walkera QR X350 Quadcopter Hovering.
https://commons.wikimedia.org/wiki/File:Walkera_QR_X350_Quadcopter_Hovering.jpg, 2013. Accessed: 03-07-22.
- [5] Markus Hehn, Robin Ritz, and Raffaello D'Andrea. Performance benchmarking of quadrotor systems using time-optimal control. *Autonomous Robots*, 33(1):69–88, 2012.

References II

- [6] Robin Ritz, Markus Hehn, Sergei Lupashin, and Raffaello D'Andrea. Quadrocopter performance benchmarking using optimal control. In *2011 IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 5179–5186. IEEE, 2011.
- [7] Zhikun Wang, Roderich Groß, and Shiyu Zhao. Aerobatic tic-toc control of planar quadcopters via reinforcement learning. *IEEE Robotics and Automation Letters*, 7(2):2140–2147, 2022.
- [8] Praveen Venkatesh, Sanket Vadhvana, and Varun Jain. Analysis and control of a planar quadrotor. *arXiv preprint arXiv:2106.15134*, 2021.
- [9] Teodor Tomić, Moritz Maier, and Sami Haddadin. Learning quadrotor maneuvers from optimal control and generalizing in real-time. In *2014 IEEE International Conference on Robotics and Automation (ICRA)*, pages 1747–1754. IEEE, 2014.

Thank you !